




Databases



Lecture 2

Structured SQL Query Language



What we know

- Data is stored according to predefined rules (data schema)
- Working with data according to predefined rules

Relational data model:

- A logical data model that does not depend on physical structures
- Based on mathematics and logic
- Relational Algebra

STRUCTURED QUERY LANGUAGE (SQL)

- Domain-specific language (Domain-specific language)
- Used to work with relational databases
- Managing a large amount of information with a single request
- No need to specify how we get the record

History

Donald Chamberlin and Ray Boyce, IBM:

- Square: (Specifying Queries As Relational Expressions)
- SEQUEL (Structured English QUery Language), 1973-1974
- Pat Selinger – cost-based optimizer
- Raymond Laurie – query compiler

Later SEQUEL -> SQL

University of California Berkeley:

- QUEL – could not stand the competition with SQL

SQL language standard

Background:

- Different software from different manufacturers
- Native query language implementation

Wanted to get:

- Software portability

Received:

- Partial portability

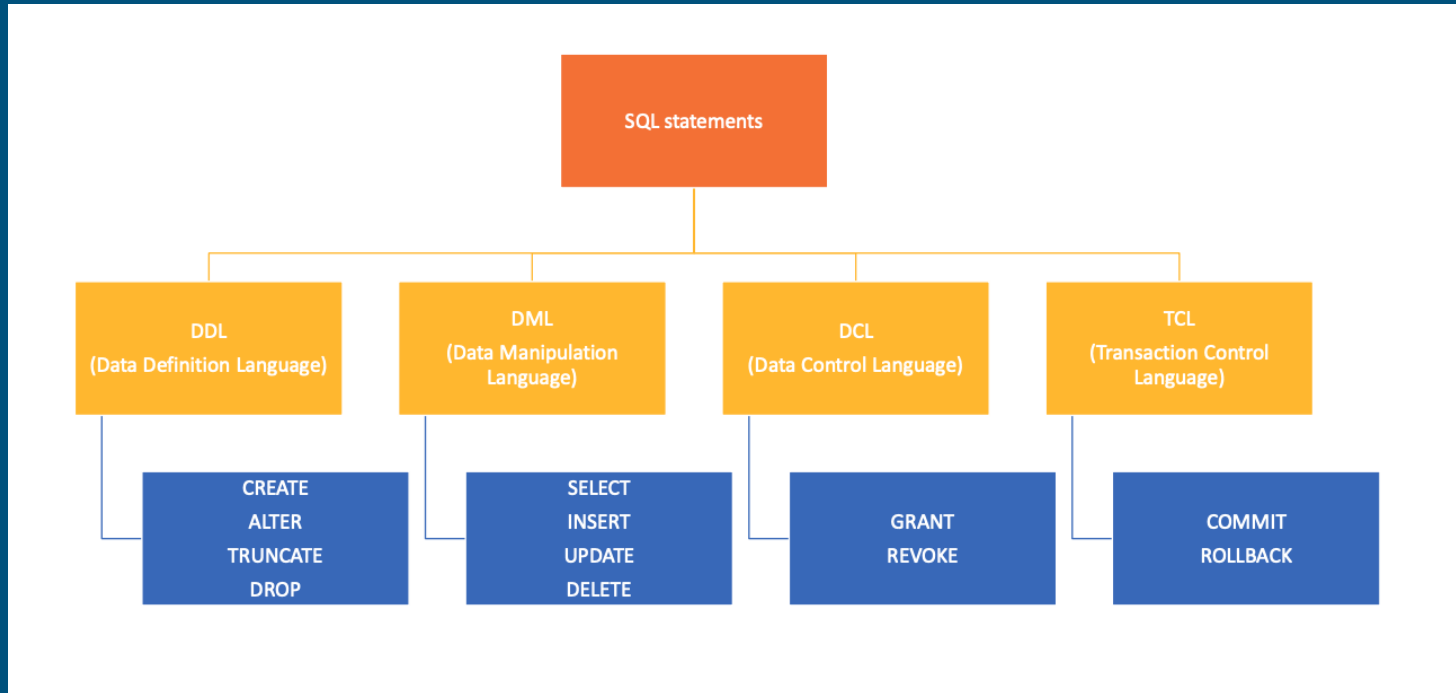
SQL language standard

1986	The first attempt at formalization
1989	SQL-86 + integrity constraint
1992	A lot of changes
1999	Regular expression matching, recursive queries, triggers, support for procedural and control operations, non-scalar types and object-oriented features.
2003	Support for SQL implementation in Java and vice versa
2006	XML-related functions, window functions, standardized sequences, and columns with automatically generated values
2008	The way SQL works with XML is defined: ways of importing and storing, publishing XML and regular data in XML format
2011	TRUNCATE, INSTEAD OF triggers. Improved window functions
2016	Adds string pattern matching, polymorphic table functions, JSON

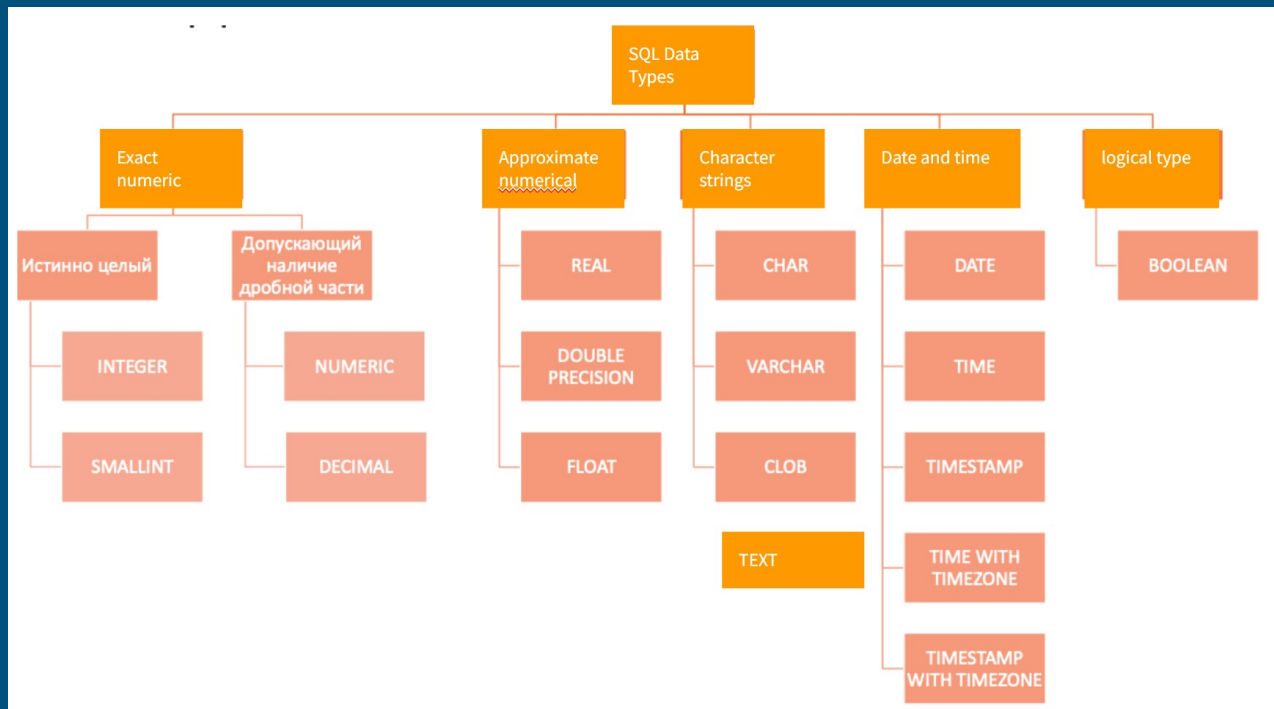
Standards issues

- Core-section of the standard (introduced in 1992)
- Manufacturers ensure compliance with Core only
- Implementation differences
- Syntax differences
- Differences in logic

SQL Statements



SQL Data Types



SQL program structure

- Program is a sequence of commands
- Command is a sequence of components. Command ends with ;
- Components are key words, identifiers, operators, strings, special symbols.
- Components are separated by space symbols.
- Each command has its own syntax.

Comments

```
CREATE TABLE lecturers (  
    id integer, --comment #1  
    name text, --comment #2  
    /* Multiline  
comments  
*/  
    birth_date date DEFAULT now() );
```

Data Definition Language

CREATE

- Database object creation operation

ALTER

- Database Object Modification Operator

DROP

- Database Object Deletion operator

TRUNCATE

- Operator for deleting the contents of a DB object

Data Definition Language

CREATE

```
CREATE [TEMPORARY] TABLE [IF NOT EXISTS] tbl_name(  
    col_name_1  datatype_1,  
    col_name_2  datatype_2,  
    ...  
    col_name_N  datatype_N  
);
```

Create

```
CREATE TABLE SUPERHERO (  
    NAME            VARCHAR(100) ,  
    AGE             INTEGER ,  
    BIRTH_DT        DATE ,  
    ALTER_EGO       VARCHAR(50)  
);
```

Create



NAME	AGE	BIRTH_DT	ALTER_EGO
------	-----	----------	-----------

Create: restrictions

CREATE TABLE students (
 student_id int PRIMARY KEY,
 first_name varchar(50),
 last_name varchar(50),
 age int NOT NULL,
 email varchar(100)
);

The diagram illustrates the components of the SQL statement. Annotations with arrows point to specific parts of the code:

- column name**: Points to `student_id`.
- column type**: Points to `int`.
- restriction**: Points to `PRIMARY KEY` and `NOT NULL`.

Other elements highlighted with green boxes include `student_id`, `int`, `PRIMARY`, and `varchar(50)` in the `last_name` definition.

Restrictions

- Serve as a means of maintaining data integrity
- They guarantee that when entering data into the table, the relationships between the data will not be broken
- Ensure referential integrity between tables (relationships) - through foreign keys
- The constraint is specified after the data type or as a separate “statement” when defining the table

Restrictions: UNIQUE

UNIQUE - a group of one or more table columns can only contain unique values:

```
CREATE TABLE products (  
  product_no integer UNIQUE,  
  name text,  
  price numeric  
);
```

```
CREATE TABLE products (  
  product_no integer,  
  name text,  
  price numeric,  
  UNIQUE (product_no)  
);
```

```
CREATE TABLE example (  
  a integer,  
  b integer,  
  c integer,  
  UNIQUE (a, c)  
);
```

NULL value

A field with a NULL value is a field with no value.

If a field in a table is not required, then you can insert a new record or update without adding a value to the field; it will be stored as NULL.

A NULL value is different from a null value (0) or a field that contains spaces (" "). It is not possible to check for NULL values with comparison operators such as =, <, or <>.

You need to use IS NULL or IS NOT NULL

Restrictions: IS NULL / IS NOT NULL

IS NULL Syntax

```
SELECT column_names  
FROM table_name  
WHERE column_name IS NULL
```

IS NOT NULL Syntax

```
SELECT column_names  
FROM table_name  
WHERE column_name IS NOT NULL
```

Restrictions: PRIMARY KEY

PRIMARY KEY (primary key) is a field (or combination fields) that uniquely identifies the record.

PRIMARY KEY should:

- contain unique values
- cannot contain NULL values.

Restrictions: PRIMARY KEY

A table cannot have two records with the same key value.

A table can only have one primary key.

```
CREATE TABLE products (  
  product_no integer UNIQUE NOT NULL,  
  name text,  
  price numeric  
);
```



```
CREATE TABLE products (  
  product_no integer PRIMARY KEY,  
  name text,  
  price numeric  
);
```

```
CREATE TABLE example (  
  a integer,  
  b integer,  
  c integer,  
  PRIMARY KEY (a, c)  
);
```

Restrictions: FOREIGN KEY

Foreign key - an attribute (or group of attributes) whose value can be repeated for several records.

Contains a reference to a primary key field in another table.

The table containing the foreign key is called a child table, containing the primary key - parent table.

Restrictions: FOREIGN KEY

```
CREATE TABLE products (  
  product_no integer PRIMARY KEY,  
  name text,  
  price numeric  
);
```



```
CREATE TABLE orders (  
  order_id integer PRIMARY KEY,  
  product_no integer REFERENCES products (product_no)  
  quantity integer  
);
```


Restrictions: FOREIGN KEY

FOREIGN KEY can refer to a group of columns of the target relation:

```
CREATE TABLE t1 (  
  a integer PRIMARY KEY,  
  b integer,  
  c integer,  
  FOREIGN KEY (b, c) REFERENCES other_table (c1, c2)  
);
```

Restrictions: CHECK

CHECK - sets an arbitrary condition on the values of one or several columns in one table row.

- The CHECK constraint specifies an expression that returns a logical result that determines whether the add or change operation is successful for specific lines:
 - TRUE or UNKNOWN - operation completed successfully
 - FALSE - an error occurs, the operation does not change anything in the database

Restrictions: CHECK

```
CREATE TABLE products (  
    product_no integer,  
    name text,  
    price numeric CHECK (price > 0)  
);
```

Restrictions: DEFAULT

The value is specified by an expression without variables (in particular, cross-references to other columns of the current table are not allowed in it). Subqueries are also not allowed. The data type of the expression that specifies the default value must match the data type of the column.

This expression will be used in all data append operations that do not specify a value for this column. If a default value is not defined, the value will be NULL.

Restrictions: DEFAULT

```
CREATE TABLE order (  
    order_id INTEGER PRIMARY KEY,  
    order_number INTEGER NOT NULL,  
    order_date  DATE DEFAULT now()::date  
);
```

Restrictions syntax

Constraints can be named using the command **CONSTRAINT** (this is useful for error output)

```
CREATE TABLE products (  
  product_no integer,  
  name text,  
  price numeric,  
  discounted_price numeric  
  CHECK (price > discounted_price)  
);
```

```
CREATE TABLE products (  
  product_no integer,  
  name text,  
  price numeric,  
  discounted_price numeric,  
  CONSTRAINT valid_discount CHECK (price >  
  discounted_price)  
);
```

Keys: potential key

A potential key is a subset of relation attributes in a relational data model that satisfies the following requirements:

- uniqueness - there are not and cannot be two tuples of a given relation in which the values of this subset of attributes coincide (are equal)
- irreducibility (minimality) - the potential key does not contain a smaller subset of attributes that satisfies the uniqueness condition. In other words, if any attribute is removed from a potential key, it will lose its uniqueness property

A candidate key always exists, even if it includes all the attributes of the relation (follows from the properties of the relations). There can be several potential keys.

Keys: primary key

Primary key (PRIMARY KEY) is such a potential key of a relationship that is selected as the “primary”

- Any potential key is suitable as a primary key

Alternate key(s) – all other potential keys of the relationship

Keys: natural key

- Natural key – a key based on an existing relation attribute
- Surrogate key – based on a special, technical attribute that has no meaningful meaning (id). Despite the obvious advantages (immutability, guaranteed uniqueness, efficiency, etc.), surrogate keys also have their disadvantages (lack of information, use instead of normalization, vulnerability of generators)

Keys: foreign key

To reflect the functional dependencies between tuples of different relations, duplication of the primary key of one relation (parent) to another (child) is used. Attributes that are copies of parent relationship keys are called foreign keys.

A foreign key in relation R2 is a non-empty subset FK of the set of attributes of this relation, such that:

- 1) There is a relation R1 with a potential key PK;
- 2) Each value of the foreign key FK in the current value of the relation R2 necessarily coincides with the value of the key PK of some tuple in the current value of the relation R1.

The relationships R1 and R2 are not necessarily different.

Ternary logic: True, False, Unknown

NOT(A)		AND(A, B)					OR(A, B)					XOR(A, B)				
		$A \wedge B$		B			$A \vee B$		B			$A \oplus B$		B		
A	$\neg A$			F	U	T			F	U	T			F	U	T
F	T		F	F	F	F		F	F	U	T		F	F	U	T
U	U	A	U	F	U	U		U	U	U	T		U	U	U	U
T	F		T	F	U	T		T	T	T	T		T	T	U	F

Ternary logic

- IS [NOT] TRUE
- IS [NOT] FALSE
- IS [NOT] UNKNOWN

p	IS TRUE	IS FALSE	IS UNKNOWN		IS NOT TRUE	IS NOT FALSE	IS NOT UNKNOWN
TRUE	TRUE	FALSE	FALSE		FALSE	TRUE	TRUE
FALSE	FALSE	TRUE	FALSE		TRUE	FALSE	TRUE
UNKNOWN	FALSE	FALSE	TRUE		TRUE	TRUE	FALSE

Data Definition Language

ALTER - modification of objects

```
ALTER TABLE table_name ADD column_name datatype;
```

```
ALTER TABLE table_name DROP column_name;
```

```
ALTER TABLE table_name RENAME column_name TO  
new_column_name;
```

```
ALTER TABLE table_name ALTER column_name TYPE datatype;
```

Alter

1. **ALTER TABLE** SUPERHERO **ADD** RATING **INTEGER**;

2. **ALTER TABLE** SUPERHERO **DROP COLUMN** AGE;

0

NAME	AGE	BIRTH_DT	ALTER_EGO
------	-----	----------	-----------

ALTER TABLE SUPERHERO **ADD** RATING **INTEGER**;



1

NAME	AGE	BIRTH_DT	ALTER_EGO	RATING
------	-----	----------	-----------	--------

ALTER TABLE SUPERHERO **DROP COLUMN** AGE;



2

NAME	BIRTH_DT	ALTER_EGO	RATING
------	----------	-----------	--------

Data Definition Language

```
TRUNCATE TABLE SUPERHERO;
```

TRUNCATE – deleting the contents of the database object (data is deleted as a whole piece, cannot be deleted by condition)

NAME	BIRTH_DT	ALTER_EGO	RATING
Tony Stark	06-JAN-1966	Iron man	100
Bruce Banner	28-FEB-1969	Hulk	80
Steve Rogers	07-MAR-1921	Captain America	90

TRUNCATE TABLE SUPERHERO;

NAME	BIRTH_DATE	ALTER_EGO	RATING
------	------------	-----------	--------

Drop

```
DROP TABLE SUPERHERO;
```

Deleting an object from the database

NAME	AGE	BIRTH_DT	ALTER_EGO
------	-----	----------	-----------



Data Manipulation Language

SELECT

- Selection of data that meets the specified conditions

INSERT

- Adding new data

UPDATE

- Change existing data

DELETE

- Delete existing data

DML: INSERT



INSERT

```
INTO table_name [(comma_separated_column_names)]  
VALUES (comma_separated_values);
```

Insert

INSERT

```
    INTO SUPERHERO (NAME, BIRTH_DT, ALTER_EGO, RATING)
VALUES ( 'Natasha Romanoff', '01-AUG-1999', 'Black Widow', 59 );
```

NAME	BIRTH_DT	ALTER_EGO	RATING
Tony Stark	06-JAN-1966	Iron man	100
Bruce Banner	28-FEB-1969	Hulk	80
Steve Rogers	07-MAR-1921	Captain America	90
Natasha Romanoff	01-AUG-1999	Black Widow	59

DML: UPDATE



```
UPDATE table_name  
    SET update_assignment_comma_list  
[WHERE conditional_expression];
```

Update

```
UPDATE SUPERHERO
    SET BIRTH_DT = '01-AUG-1940'
    WHERE NAME = 'Natasha Romanoff';
```


NAME	BIRTH_DATE	ALTER_EGO	RATING
Tony Stark	06-JAN-1966	Iron man	100
Bruce Banner	28-FEB-1969	Hulk	80
Steve Rogers	07-MAR-1921	Captain America	90
Natasha Romanoff	01-AUG-1940	Black Widow	59

DML: DELETE

DELETE

FROM table_name

[**WHERE** conditional_expression];



Delete

DELETE

FROM SUPERHERO

WHERE NAME = 'Bruce Banner';

NAME	BIRTH_DT	ALTER_EGO	RATING
Tony Stark	06-JAN-1966	Iron man	100
Bruce Banner	28-FEB-1969	Hulk	80
Steve Rogers	07-MAR-1921	Captain America	90
Natasha Romanoff	01-AUG-1940	Black Widow	59

DML: DELETE vs DDL:TRUNCATE


Delete	Truncate
"Delete" line by line	Deleting the entire block at once
You can set conditions for deletion	Conditions for deletion cannot be set
Ability to roll back changes	There is no possibility of rolling back changes
Physically, the lines are not deleted, only marked "invisible" from a certain moment VACUUM is needed for removal	Deleting data and freeing up disk space occurs immediately

DML: SELECT



```
SELECT [DISTINCT] select_item_comma_list  
  FROM table_reference_comma_list  
[WHERE conditional_expression]  
[GROUP BY column_name_comma_list]  
[HAVING conditional_expression]  
[ORDER BY order_item_comma_list];
```

The order of execution of the request



SELECT: FROM

```
SELECT ALTER_EGO  
FROM SUPERHERO;
```

```
SELECT NAME, ALTER_EGO, COMICS_N  
FROM SUPERHERO, COMICS;
```

NAME	BIRTH_DT	ALTER_EGO	RATING
Tony Stark	06-JAN-1966	Iron man	98
Bruce Banner	28-FEB-1969	Hulk	80
Steve Rogers	07-MAR-1921	Captain America	90
Natasha Romanoff	01-AUG-1940	Black Widow	59
Thor	13-FEB-1976	Thor	74
Clint Barton	17-DEC-1969	Hawkeye	55
Wanda Maximoff	22-OCT-1974	Scarlet Witch	81
Pietro Maximoff	22-OCT-1974	Quicksilver	82
Charles Xavier	30-JUN-1933	Professor X	100
Jean Grey	12-SEP-1961	Phoenix	93
Wade Wilson	13-APR-1980	Deadpool	89
James Howlett	01-JAN-1887	Wolverine	99

SELECT: WHERE

```
SELECT NAME, ALTER_EGO
FROM SUPERHERO
WHERE RATING > 90;
```

NAME	ALTER_EGO
Tony Stark	Iron Man
Charles Xavier	Professor X
Jean Grey	Phoenix
James Howlett	Wolverine

```
SELECT NAME, ALTER_EGO
FROM SUPERHERO
WHERE RATING < 50;
```

NAME	ALTER_EGO

Where

- **WHERE** X = value_1 **AND** X <> value_2;
- **WHERE** X = value_1 **OR** X <> value_2;
- **WHERE** X = value_1 **AND NOT** X < value_3;
- **WHERE** X < value_1 **AND** X > value_2 **OR** X = value_3

priority: not, and, or

Useful functions

Sometimes it is useful to use special functions in a request:

IN - belonging to a certain set of values:

$X \text{ IN } (a_1, a_2, \dots, a_n) \equiv X = a_1 \text{ or } X = a_2 \text{ or } \dots \text{ or } X = a_n$

BETWEEN - belonging to a certain range of values:

$X \text{ BETWEEN } A \text{ AND } B \equiv (X \geq A \text{ and } X \leq B) \text{ or } (X \leq A \text{ and } X \geq B)$

LIKE - satisfying the text pattern:

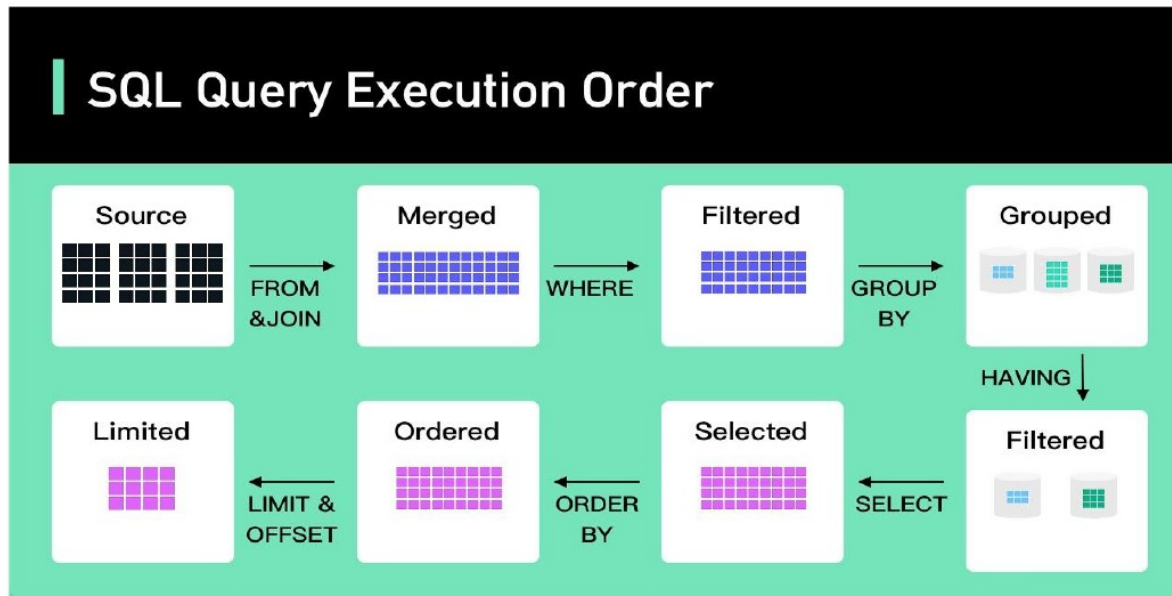
$X \text{ LIKE } '0\%abc_0'$, where $_$ is exactly 1 character, and $\%$ is any sequence of characters (including zero length).

```
SELECT DATE BIRTH
        , RATING
        , NAME
FROM SUPERHERO
WHERE RATING < 90
        AND RATING > 70
GROUP BY DATE_BIRTH;
```

Why it would not work?

Steps

1. FROM / JOIN
2. WHERE
3. GROUP BY
4. HAVING
5. SELECT
6. ORDER BY
7. LIMIT / OFFSET



Aggregating functions

- `count()` - the number of records with a known value. If you need to count the number of unique values, you can use `count(DISTINCT field_nm)`
- `max()` - the largest of all selected field values
- `min()` - the smallest of all selected field values
- `sum()` - sum of all selected field values
- `avg()` - average of all selected field values

Count

```
SELECT count (ALTER_EGO)
FROM SUPERHERO;
```

COUNT (ALTER_EGO)

12

NAME	BIRTH_DT	ALTER_EGO	RATING
Tony Stark	06-JAN-1966	Iron man	98
Bruce Banner	28-FEB-1969	Hulk	80
Steve Rogers	07-MAR-1921	Captain America	90
Natasha Romanoff	01-AUG-1940	Black Widow	59
Thor	13-FEB-1976	Thor	74
Clint Barton	17-DEC-1969	Hawkeye	55
Wanda Maximoff	22-OCT-1974	Scarlet Witch	81
Pietro Maximoff	22-OCT-1974	Quicksilver	82
Charles Xavier	30-JUN-1933	Professor X	100
Jean Grey	12-SEP-1961	Phoenix	93
Wade Wilson	13-APR-1980	Deadpool	89
James Howlett	01-JAN-1887	Wolverine	99

Group by

```
SELECT BIRTH_DT,  
        count(ALTER_EGO)  
FROM SUPERHERO  
WHERE BIRTH_DT = '22-OCT-1974'  
GROUP BY BIRTH_DT;
```

BIRTH_DT	COUNT (ALTER_EGO)
22-OCT-1974	2

Having

- It is used in conjunction with GROUP BY to impose restrictions on the selection after grouping
- The restriction with the use of WHERE can be imposed only before grouping

```
GROUP BY column_name(s)  
HAVING expression_clause
```

```
SELECT ITEM,  
        avg(PRICE)  
FROM CATALOG  
GROUP BY ITEM;
```

ITEM	AVG (PRICE)
Computer	1035.67
Laptop	1000
Mobile phone	200
Printer	300
Scanner	200
Camera	525
Headphones	200

```
SELECT ITEM, avg(PRICE)
FROM CATALOG
GROUP BY ITEM
HAVING avg(PRICE) <= 500;
```

ITEM	AVG (PRICE)
Mobile phone	200
Printer	300
Scanner	200
Headphones	200

Order by

```
SELECT ID, ITEM, MAGAZINE, PRICE, DELIVERY  
      FROM CATALOG  
      ORDER BY PRICE ASC;
```

ASC – ascending

DESC – descending

Distinct

```
SELECT DISTINCT  
    ITEM  
FROM CATALOG;
```

ID	ITEM	MAGAZINE	PRICE	DELIVERY
5	Mobile phone	1	150	15
9	Scanner	2	200	7
12	Headphones	4	200	0
7	Mobile phone	3	250	10
8	Printer	1	300	15
10	Camera	2	500	5
11	Camera	3	550	10
4	Laptop	1	999	15
1	Computer	1	1000	15
6	Laptop	2	1001	5
2	Computer	2	1007	5
3	Computer	3	1100	10

ITEM
Computer
Laptop
Mobile phone
Printer
Scanner
Camera
Headphones

Select *

```
SELECT *  
  FROM CATALOG  
 WHERE ITEM = 'Scanner';
```

ID	ITEM	MAGAZINE	PRICE	DELIVERY
9	Scanner	2	200	7

Select: alias

```
SELECT column_name AS alias_column_name  
FROM table_name alias_table_name;
```

```
SELECT MAGAZINE,  
        count(ITEM) AS cnt_items  
FROM CATALOG  
WHERE DELIVERY < 15  
GROUP BY MAGAZINE  
HAVING count(ITEM) > 1  
ORDER BY cnt_items;
```

More aliases

Alias Column Syntax

```
SELECT column_name AS alias_name  
FROM table_name;
```

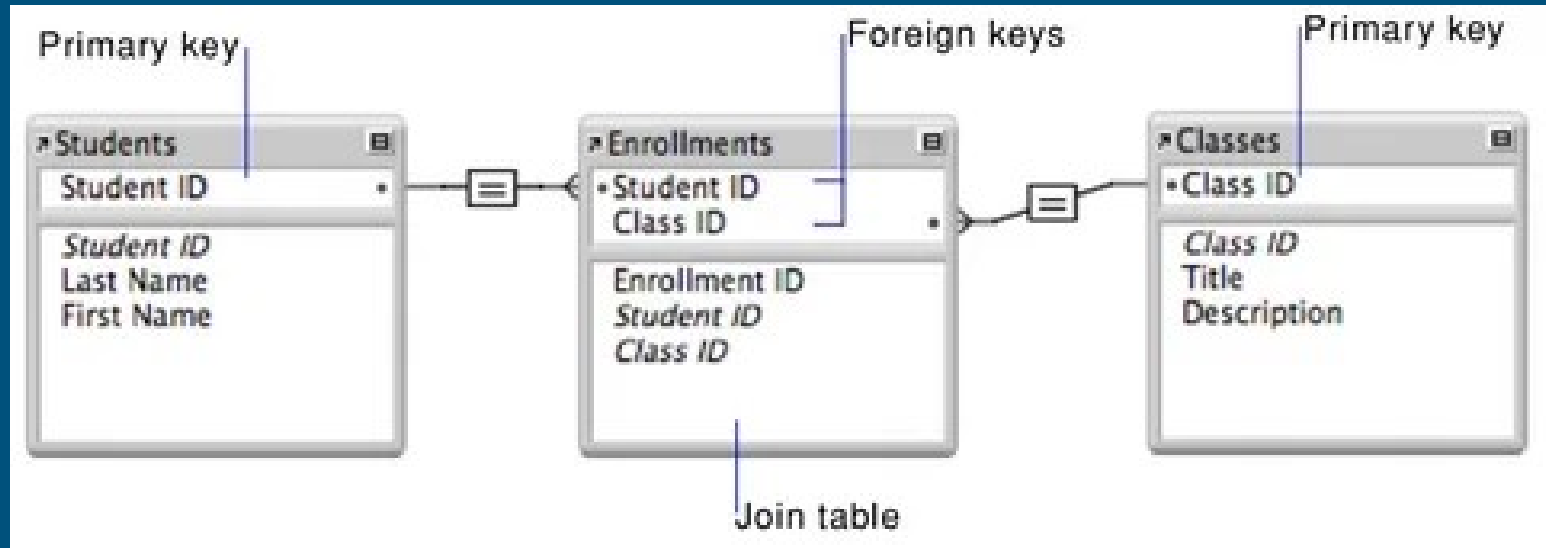
Alias Table Syntax

```
SELECT column_name(s)  
FROM table_name AS alias_name;
```

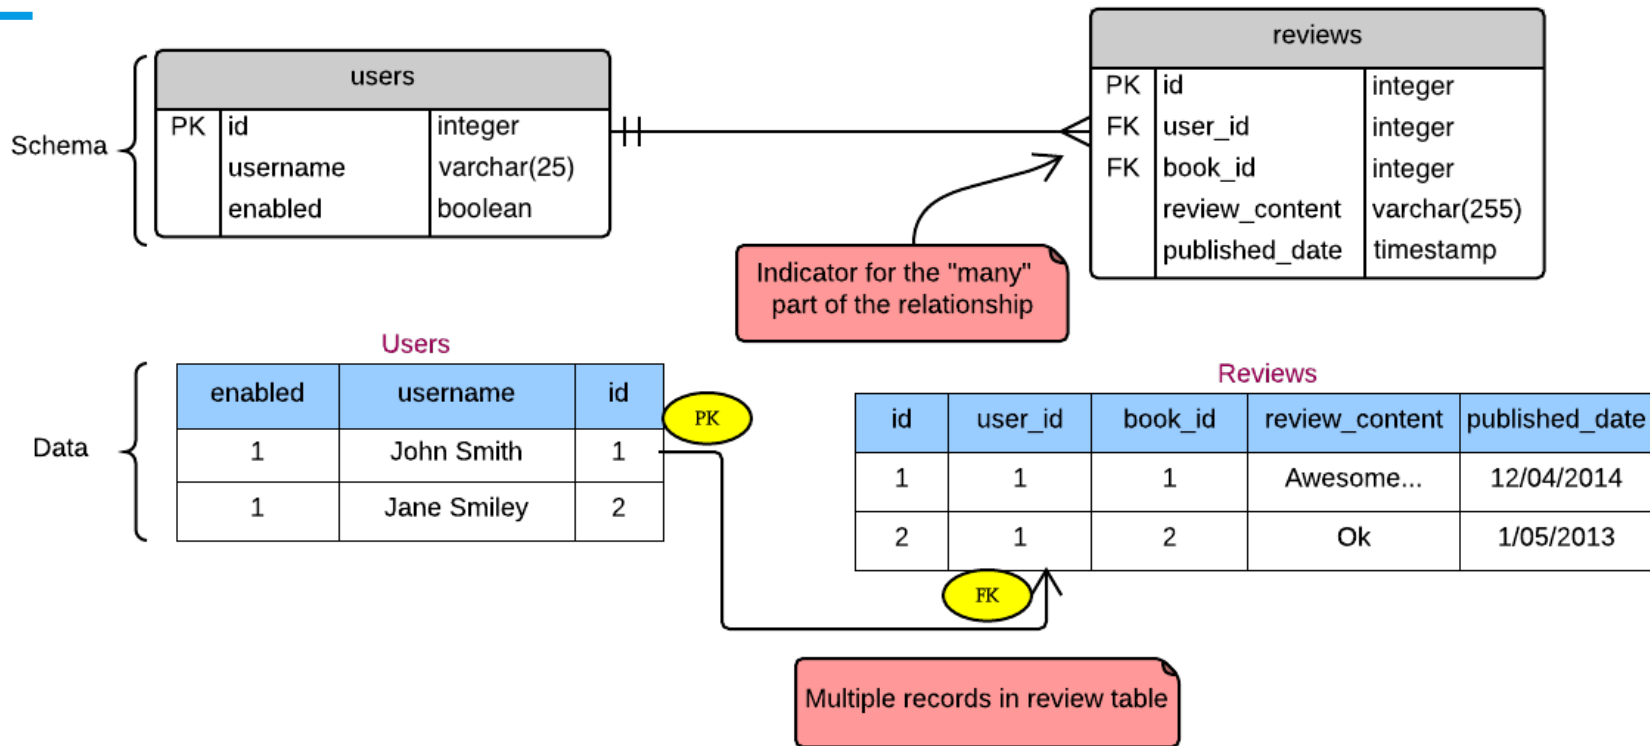
Relations between tables

- one to one
- one to many
- many to many

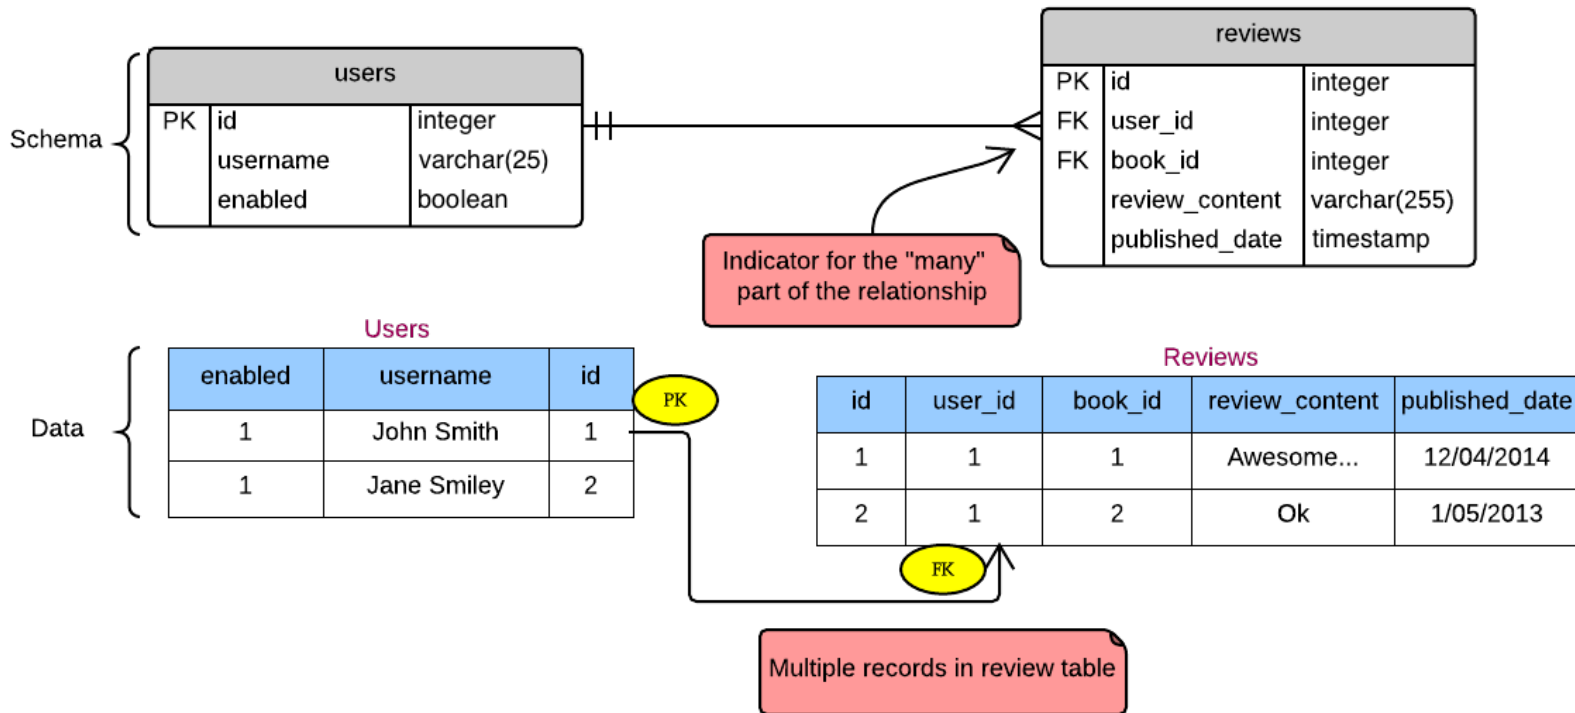
FOREIGN KEY



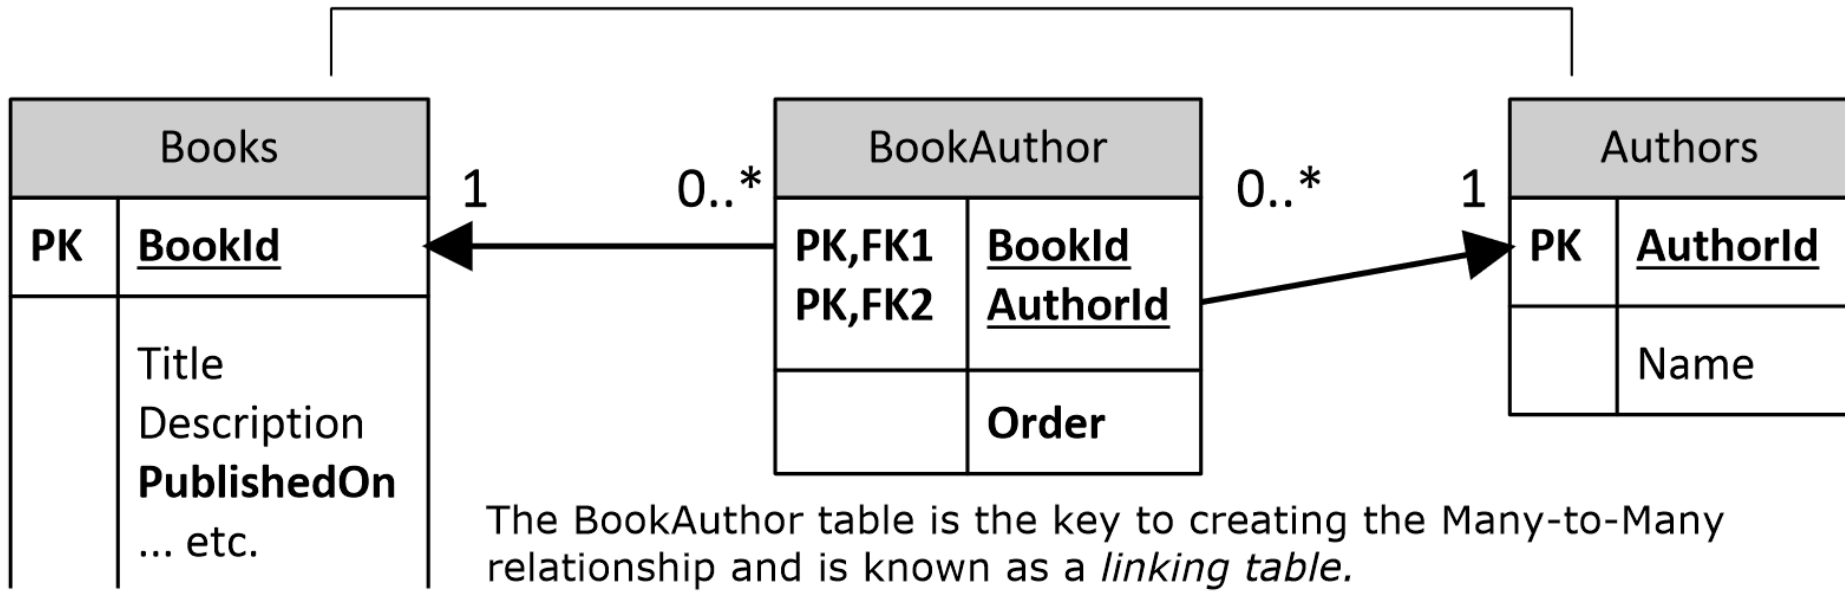
One-to-one




One-to-many



Many-to-many



JOIN



INNER JOIN

INNER JOIN: Returns records that have matching values in both tables

INNER JOIN Syntax

```
SELECT column_name(s)  
FROM table1  
INNER JOIN table2 ON table1.column_name = table2.column_name
```

INNER JOIN

Customers

CustomerId	Name
1	Shree
2	Kalpna
3	Basavaraj

Orders

OrderId	CustomerId	OrderDate
100	1	2014-01-29 23:56:57.700
200	4	2014-01-30 23:56:57.700
300	3	2014-01-31 23:56:57.700

**INNER JOIN on CustomerId
Column**

RESULT

CustomerId	Name	OrderId	CustomerId	OrderDate
1	Shree	100	1	2014-01-30 23:48:32.850
3	Basavaraj	300	3	2014-02-01 23:48:32.853

LEFT JOIN

LEFT (OUTER) JOIN: Returns all records from the left table and matching records from the table on the right

LEFT JOIN Syntax

```
SELECT column_name(s)  
FROM table1  
LEFT JOIN table2 ON table1.column_name = table2.column_name;
```

LEFT OUTER JOIN

Customers

CustomerId	Name
1	Shree
2	Kalpana
3	Basavaraj

Orders

OrderId	CustomerId	OrderDate
100	1	2014-01-29 23:56:57.700
200	4	2014-01-30 23:56:57.700
300	3	2014-01-31 23:56:57.700

**LEFT OUTER JOIN on
CustomerId Column**

RESULT

CustomerId	Name	OrderId	CustomerId	OrderDate
1	Shree	100	1	2014-01-30 23:48:32.850
2	Kalpana	NULL	NULL	NULL
3	Basavaraj	300	3	2014-02-01 23:48:32.853

RIGHT JOIN

RIGHT (OUTER) JOIN: Returns all records from the table on the right, and matching records from the left table

RIGHT JOIN Syntax

```
SELECT column_name(s)  
FROM table1  
RIGHT JOIN table2 ON table1.column_name = table2.column_name;
```

RIGHT OUTER JOIN

Customers

CustomerId	Name
1	Robert
2	Peter
3	Smith

Orders

OrderId	CustomerId	OrderDate
100	1	2016-10-19 15:21:27
200	4	2016-10-20 15:21:27
300	2	2016-10-21 15:21:27

**RIGHT OUTER JOIN on
CustomerId Column**

RESULT

CustomerId	Name	OrderId	CustomerId	OrderDate
1	Robert	100	1	2016-10-19 15:21:27
NULL	NULL	200	4	2016-10-20 15:21:27
2	Peter	300	2	2016-10-21 15:21:27

FULL JOIN

FULL (OUTER) JOIN: Returns all records when there is a match in the left or right table

FULL OUTER JOIN Syntax

```
SELECT column_name(s)  
FROM table1  
FULL OUTER JOIN table2 ON table1.column_name = table2.column_name;
```

FULL OUTER JOIN

Customers

CustomerId	Name
1	Shree
2	Kalpana
3	Basavaraj

Orders

OrderId	CustomerId	OrderDate
100	1	2014-01-29 23:56:57.700
200	4	2014-01-30 23:56:57.700
300	3	2014-01-31 23:56:57.700

**FULL OUTER JOIN on
CustomerId Column**

RESULT

CustomerId	Name	OrderId	CustomerId	OrderDate
1	Shree	100	1	2014-01-30 23:48:32.850
2	Kalpana	NULL	NULL	NULL
3	Basavaraj	300	3	2014-02-01 23:48:32.853
NULL	NULL	200	4	2014-01-31 23:48:32.853

CROSS JOIN

CROSS JOIN = Cartesian product of two tables

```
SELECT column_name(s)  
FROM table1  
CROSS JOIN table2
```


CROSS JOIN

Customers

CustomerId	Name
1	Shree
2	Kalpana
3	Basavaraj

Orders

OrderId	CustomerId	OrderDate
100	1	2014-01-29 23:56:57.700
200	4	2014-01-30 23:56:57.700
300	3	2014-01-31 23:56:57.700




CustomerId	Name	OrderId	CustomerId	OrderDate
1	Shree	100	1	2014-01-30 23:48:32.850
2	Kalpana	100	1	2014-01-30 23:48:32.850
3	Basavaraj	100	1	2014-01-30 23:48:32.850
1	Shree	200	4	2014-01-31 23:48:32.853
2	Kalpana	200	4	2014-01-31 23:48:32.853
3	Basavaraj	200	4	2014-01-31 23:48:32.853
1	Shree	300	3	2014-02-01 23:48:32.853
2	Kalpana	300	3	2014-02-01 23:48:32.853
3	Basavaraj	300	3	2014-02-01 23:48:32.853

Subquery example

```
PROC SQL ;
CREATE TABLE prim2 AS
  SELECT pt_id, admdate, disdate, hosp, md_id,
         b.lastname AS ptname,
         c.lastname AS mdname
  FROM ex.admits a, ex.patients b,
       (SELECT DISTINCT md_id, lastname
        FROM ex.doctors) c
  WHERE (a.pt_id EQ b.id) AND
        (a.md EQ b.primmd) AND
        (a.md EQ c.md_id)
  ORDER BY a.pt_id, admdate ;
QUIT;
```

Create inline table C to join
with A and B using IDs



Subquery

Subqueries (also known as inner queries or nested queries) are a tool for performing operations in multiple steps. For example, if you wanted to take the sums of several columns, then average all of those values, you'd need to do each aggregation in a distinct step.

Subquery example 1

```
SELECT sub.*  
FROM (  
    SELECT *  
    FROM tutorial.sf_crime_incidents_2014_01  
    WHERE day_of_week = 'Friday'  
) sub  
WHERE sub.resolution = 'NONE'
```

Subquery example 2

```
SELECT *  
  FROM tutorial.sf_crime_incidents_2014_01  
 WHERE Date = (SELECT MIN(date)  
               FROM tutorial.sf_crime_incidents_2014_01  
              )
```


Subquery example 3

```
SELECT *
```

```
FROM tutorial.sf_crime_incidents_2014_01
```

```
WHERE Date IN (SELECT date
```

```
FROM tutorial.sf_crime_incidents_2014_01
```

```
ORDER BY date
```

```
LIMIT 5
```

```
)
```

Subquery example 4

```
SELECT incidents.*, sub.incidents AS incidents_that_day
FROM tutorial.sf_crime_incidents_2014_01 incidents
JOIN ( SELECT date, COUNT(incident_num) AS incidents
      FROM tutorial.sf_crime_incidents_2014_01
      GROUP BY 1 ) sub
ON incidents.date = sub.date
ORDER BY sub.incidents DESC, time
```

UNIONS

```
SELECT *
```

```
FROM  
tutorial.crunchbase_investments_par  
t1
```

```
UNION ALL
```

```
SELECT *
```

```
FROM  
tutorial.crunchbase_investments_par  
t2;
```

SQL has strict rules for
appending data:

1. Both tables must have the same number of columns
2. The columns must have the same data types in the same order as the first table